

metrics_analysis_report

April 13, 2025

1 People Metrics & Retention Insights

(Using synthetic data)

1.1 Importing Libraries

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Modeling
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report

# Seaborn Plotting theme
sns.set_theme(style='whitegrid')

print("Libraries imported successfully.")
```

Libraries imported successfully.

1.2 Importing Dataset

```
[2]: df = pd.read_csv("data/employee_metrics.csv")

print("Dataset imported successfully.")
```

Dataset imported successfully.

1.3 Viewing the Dataset

```
[3]: df.head()
```

```
[3]:   EmployeeID  JoinDate LastPromotionDate  Age Gender Department Location \
0           1  2020-06-15      2021-06-15   32  Male          HR      Berlin
1           2  2015-05-11      2024-05-08   50  Male          Sales    Berlin
2           3  2016-06-01      2024-05-30   28  Male          Product  Berlin
```

3	4	2012-09-19	2013-09-19	50	Other	Product	Berlin
4	5	2019-06-16	2019-06-16	33	Male	Marketing	Dubai

	JobLevel	EngagementScore	PerformanceScore	ManagerSatisfaction	\
0	4	4.276240	3.818902	3.834850	
1	4	8.521031	2.919122	2.869796	
2	1	6.983925	1.586720	1.430099	
3	4	5.375425	3.071954	2.609112	
4	3	9.431298	3.274224	3.581034	

	NumTrainingsCompleted	TenureYears	LeftCompany	AttritionReason
0	4	4	0	NaN
1	8	9	1	Retired
2	10	8	0	NaN
3	1	12	0	NaN
4	4	5	0	NaN

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   EmployeeID            5000 non-null   int64
1   JoinDate              5000 non-null   object
2   LastPromotionDate     5000 non-null   object
3   Age                   5000 non-null   int64
4   Gender                5000 non-null   object
5   Department            5000 non-null   object
6   Location              5000 non-null   object
7   JobLevel              5000 non-null   int64
8   EngagementScore       5000 non-null   float64
9   PerformanceScore      5000 non-null   float64
10  ManagerSatisfaction    5000 non-null   float64
11  NumTrainingsCompleted  5000 non-null   int64
12  TenureYears           5000 non-null   int64
13  LeftCompany           5000 non-null   int64
14  AttritionReason       1502 non-null   object
dtypes: float64(3), int64(6), object(6)
memory usage: 586.1+ KB
```

```
[5]: df.describe()
```

	EmployeeID	Age	JobLevel	EngagementScore	PerformanceScore	\
count	5000.000000	5000.0000	5000.000000	5000.000000	5000.000000	
mean	2500.500000	38.5792	3.006000	6.472303	3.492530	
std	1443.520003	9.6542	1.414908	1.961729	0.940692	

min	1.000000	22.0000	1.000000	1.000000	1.000000
25%	1250.750000	30.0000	2.000000	5.144574	2.853289
50%	2500.500000	39.0000	3.000000	6.518082	3.519815
75%	3750.250000	47.0000	4.000000	7.892489	4.192822
max	5000.000000	55.0000	5.000000	10.000000	5.000000

	ManagerSatisfaction	NumTrainingsCompleted	TenureYears	LeftCompany
count	5000.000000	5000.000000	5000.000000	5000.000000
mean	3.463982	4.986400	7.253200	0.300400
std	1.058885	3.166231	3.209914	0.458478
min	1.000000	0.000000	2.000000	0.000000
25%	2.716705	2.000000	4.000000	0.000000
50%	3.522048	5.000000	7.000000	0.000000
75%	4.319385	8.000000	10.000000	1.000000
max	5.000000	10.000000	13.000000	1.000000

1.4 Transformation

```
[6]: # Changing Dates to datetime format
df['JoinDate'] = pd.to_datetime(df['JoinDate'])
df['LastPromotionDate'] = pd.to_datetime(df['LastPromotionDate'])

print('Datetime formatted.')
```

Datetime formatted.

```
[7]: # Calculating years since promotion for analysis
df["YearsSincePromotion"] = ((pd.to_datetime("2025-04-01") -
    ↪df["LastPromotionDate"]).dt.days // 365) # Calculating how many years since
    ↪promotion from 2025-04-01

df.head()
```

```
[7]: EmployeeID  JoinDate  LastPromotionDate  Age  Gender  Department  Location  \
0           1  2020-06-15      2021-06-15   32   Male      HR      Berlin
1           2  2015-05-11      2024-05-08   50   Male      Sales     Berlin
2           3  2016-06-01      2024-05-30   28   Male     Product     Berlin
3           4  2012-09-19      2013-09-19   50  Other     Product     Berlin
4           5  2019-06-16      2019-06-16   33   Male    Marketing    Dubai
```

```
JobLevel  EngagementScore  PerformanceScore  ManagerSatisfaction  \
0           4           4.276240           3.818902           3.834850
1           4           8.521031           2.919122           2.869796
2           1           6.983925           1.586720           1.430099
3           4           5.375425           3.071954           2.609112
4           3           9.431298           3.274224           3.581034
```

```
NumTrainingsCompleted  TenureYears  LeftCompany  AttritionReason  \
```

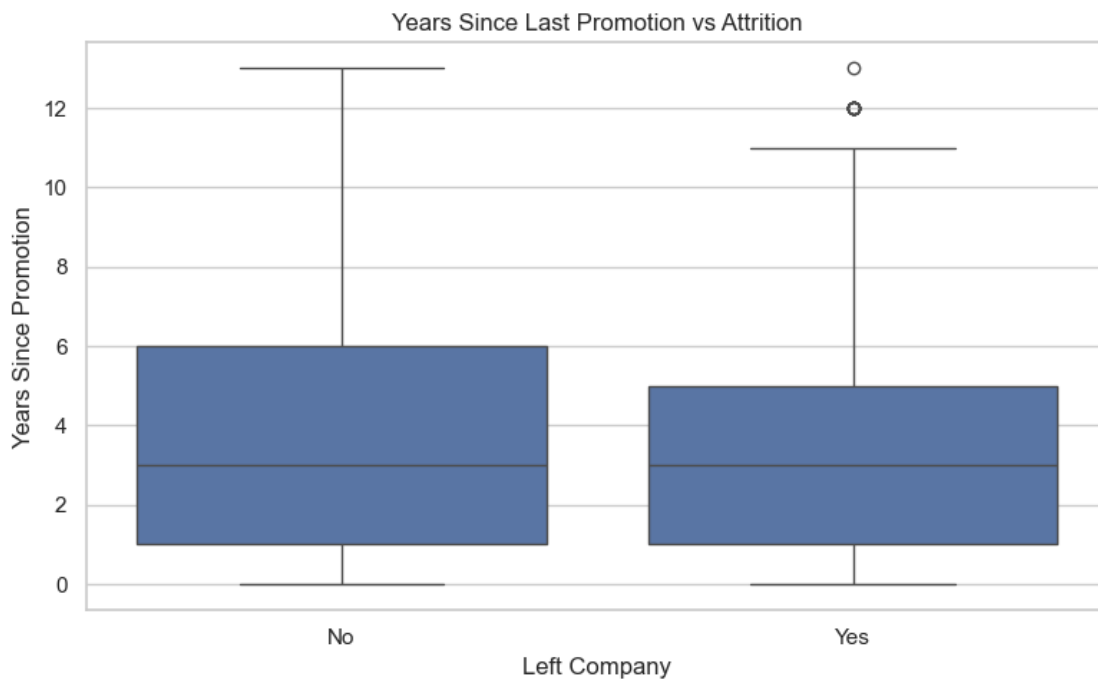
0	4	4	0	NaN
1	8	9	1	Retired
2	10	8	0	NaN
3	1	12	0	NaN
4	4	5	0	NaN

	YearsSincePromotion
0	3
1	0
2	0
3	11
4	5

1.5 Analysis

1.5.1 Promotion vs Retention Analysis

```
[8]: plt.figure(figsize=(8, 5))
sns.boxplot(x="LeftCompany", y="YearsSincePromotion", data=df)
plt.title("Years Since Last Promotion vs Attrition")
plt.xticks(ticks=[0, 1], labels=["No", "Yes"])
plt.xlabel("Left Company")
plt.ylabel("Years Since Promotion")
plt.tight_layout()
plt.show()
```

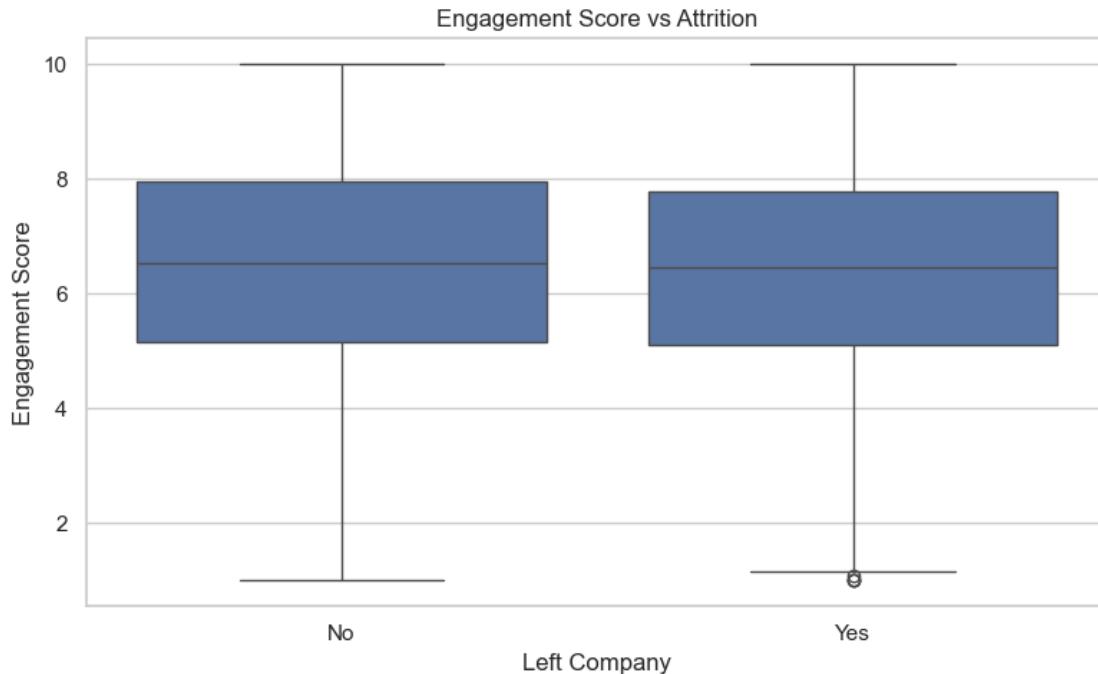


The similar distributions suggest that time since last promotion isn't a strong differentiator between employees who stay versus those who leave.

However, the presence of outliers specifically in the "Yes" group indicates that employees who haven't been promoted for exceptionally long periods (12+ years) may eventually choose to leave the company.

1.6 Engagement vs Attrition

```
[9]: plt.figure(figsize=(8, 5))
sns.boxplot(x="LeftCompany", y="EngagementScore", data=df)
plt.title("Engagement Score vs Attrition")
plt.xlabel("Left Company")
plt.xticks(ticks=[0, 1], labels=["No", "Yes"])
plt.ylabel("Engagement Score")
plt.tight_layout()
plt.show()
```



Surprisingly, engagement scores appear to be almost identical between employees who left and those who stayed, suggesting that employee attrition at this company may be driven by factors other than engagement levels.

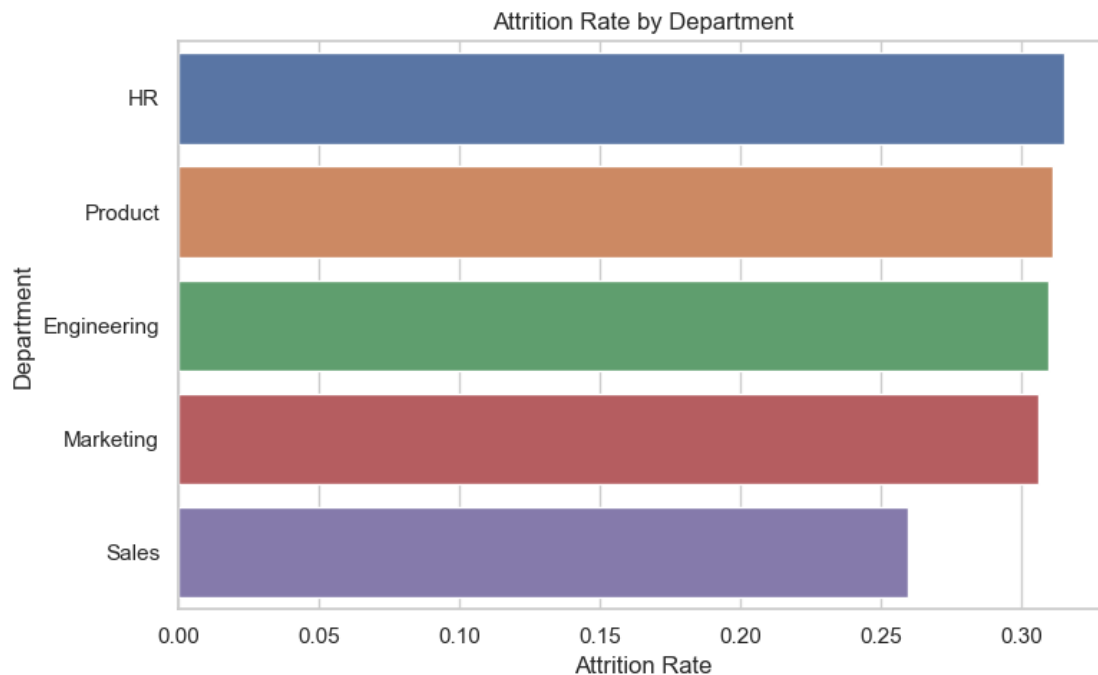
This challenges the common assumption that disengaged employees are more likely to leave, indicating management should look beyond engagement metrics to understand and address turnover issues.

1.7 Attrition Rate by Department

```
[10]: dep_attrition = df.groupby("Department")["LeftCompany"].mean().  
      ↪sort_values(ascending=False)  
  
      dep_attrition.head()
```

```
[10]: Department  
HR          0.315377  
Product     0.311245  
Engineering  0.309524  
Marketing    0.306101  
Sales        0.259921  
Name: LeftCompany, dtype: float64
```

```
[11]: plt.figure(figsize=(8, 5))  
      sns.barplot(x=dep_attrition, y=dep_attrition.index, hue=dep_attrition.index)  
      plt.title("Attrition Rate by Department")  
      plt.xlabel("Attrition Rate")  
      plt.tight_layout()  
      plt.show()
```



The technical and administrative departments (HR, Product, and Engineering) are experiencing significantly higher turnover than Sales.

With attrition rates approaching or exceeding 30% in four out of five departments, the

company is facing a substantial retention problem across most of the organization.

The notably lower turnover in Sales suggests this department may have more effective retention strategies or compensation structures that could potentially be applied elsewhere in the company to address the widespread retention challenges.

1.8 Prediction Modeling

1.8.1 Logistic Regression Model

```
[12]: # Select Features
features = ["EngagementScore", "PerformanceScore", "YearsSincePromotion",
           ↪ "TenureYears", "ManagerSatisfaction"]

# Assign X and y
X = df[features]
y = df["LeftCompany"]
```

```
[13]: # Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
           ↪ random_state=42)
```

```
[14]: # Train Model
model = LogisticRegression(class_weight='balanced')
model.fit(X_train, y_train)
```

```
[14]: LogisticRegression(class_weight='balanced')
```

1.8.2 Prediction

```
[15]: y_pred = model.predict(X_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.70	0.50	0.58	685
1	0.33	0.53	0.41	315
accuracy			0.51	1000
macro avg	0.51	0.52	0.50	1000
weighted avg	0.58	0.51	0.53	1000

1.9 Project Insights and Recommendations

1.9.1 Key Recommendations

1. Address Long-Term Promotion Stagnation

- Target employees with 8+ years since last promotion as they represent significant attrition risk, despite overall promotion timing not being a strong predictor.

- Implement career pathing reviews for employees approaching 5 years without advancement.
2. **Look Beyond Traditional Engagement Metrics**
 - Engagement scores show minimal correlation with attrition decisions, suggesting deeper factors at play.
 - Conduct targeted exit interviews to identify true retention drivers beyond standard engagement surveys.
 3. **Department-Specific Retention Strategies**
 - Prioritize HR department (32% attrition) for immediate intervention.
 - Implement tailored retention programs for Product and Engineering departments (30%+ attrition).
 - Study Sales practices (25% attrition) to identify transferable retention success factors.
 4. **Comprehensive Retention Framework**
 - Develop a holistic approach combining targeted promotions, growth opportunities, and department-specific initiatives.
 - Establish quarterly retention risk assessments focusing on employees with multiple risk factors.